A Dual-Expert Pipeline Combining Classical Image Processing, Biomarker Machine Learning, MobileNetV3 and Temporal Bi-LSTM for Diabetic Retinopathy Severity Analysis in Fundus Images and Video

Group N053, N149 and N150
Department of Computer Engineering, NMIMS University, Mumbai, India
Image and Video Processing (IVP) Project

Abstract*—Diabetic retinopathy (DR) is a leading cause of preventable adult blindness globally, yet timely screening is constrained by grader shortages and variable fundus image quality. This paper presents an end-to-end, CPU-optimized diagnostic system integrating classical retinal image processing, quantitative biomarker extraction, deep transfer learning, and temporal sequence modeling for five-grade DR severity analysis on both static fundus photographs and continuous retinal video stream sweeps (.mp4). The pipeline executes Ben Graham circular pupil cropping, green-channel isolation, and CLAHE normalization. A morphological feature engine extracts twelve quantitative clinical biomarkers, including Frangi Hessian vesselness, a 6-angle Gabor wavelet bank, optic disc diameter, hard exudates, microaneurysm spot counts, and foveal proximity for Clinically Significant Macular Edema (CSME) detection. Diagnostic inference is driven by a ***Dual-Expert Hybrid Consensus Architecture fusing deep spatial representations (MobileNetV3-Small with 2x Test-Time Augmentation) with an ensemble of classical machine learning classifiers (Random Forest + HistGradientBoosting) trained across 750 balanced patient images. For video scans, frames undergo real-time Laplacian focus gating (Var >= 15.0), ORB feature registration, and RANSAC homography to reconstruct a noise-free super-resolution median composite, alongside a dedicated 16-frame Video CNN + Bi-LSTM with Temporal Attention (RetinalVideoCNN_LSTM). Evaluated on balanced clinical benchmarks, the system achieves ***92.50% Primary DR Screening Accuracy*** on photos and ***96.00% Binary Screening Accuracy*** on video sweeps (exceeding the >=90% clinical safety standard). For 5-class severity grading, the pipeline delivers ***80.00% Clinical Adjacent-Grade Agreement (+-1 stage)***, ***62.50% Exact 5-Class Match*** (80% macro precision, 100% precision on Grade 2 and 4), and a Quadratic Weighted Kappa (QWK) of ***0.5590*** on photos and ***0.5287*** on video. CPU thread throttling, 4x sub-sampled Frangi filtering, and an O(1) video streaming memory architecture (12x RAM reduction) ensure lag-free CPU execution. The system is deployed as an interactive Streamlit clinical dashboard with automated PDF and HTML diagnostic report generation.*

Keywords—Diabetic retinopathy, fundus image processing, Frangi filter, Gabor wavelets, MobileNetV3, Bi-LSTM, Grad-CAM, Dual-Expert consensus, retinal video analysis.

I. Introduction

Diabetic retinopathy (DR) is a severe microvascular complication of diabetes leading to progressive visual impairment and blindness. A global systematic meta-analysis estimated that 103.12 million adults had DR in 2020 and projected 160.50 million by 2045, with a pooled prevalence of 22.27% among diabetic individuals [1]. Early DR is frequently asymptomatic, making mass screening vital. Disease severity conforms to the International Clinical Diabetic Retinopathy (ICDR) scale: Grade 0 (No DR), Grade 1 (Mild NPDR - microaneurysms only), Grade 2 (Moderate NPDR), Grade 3 (Severe NPDR - 4-2-1 rule), and Grade 4 (Proliferative DR - neovascularization) [2]. Automated mass screening faces critical hurdles: (i) illumination and focus vary widely across clinics; (ii) pure deep learning models act as uninterpretable black boxes whose saliency maps do not guarantee lesion fidelity [7], [8]; (iii) high-end GPU hardware is rarely available in primary healthcare clinics; and (iv) handheld ophthalmoscopes increasingly capture video sweeps rather than static photographs. This work presents an integrated, CPU-optimized, Dual-Expert diagnostic system that resolves these challenges by coupling digital image processing, physical biomarker extraction, deep transfer learning, and temporal video sequence modeling.

II. Related Work & Engineering Inferences

Classical retinal processing isolates the green channel for maximum hemoglobin absorption contrast and employs CLAHE [11] and Ben Graham local subtraction [12] for illumination equalization. Vessel enhancement relies on multiscale Frangi Hessian filtering [13] and Gabor wavelet banks [14]. Microaneurysm detection employs morphological top-hat filtering and candidate classification [16], [17]. For tabular classification on physical biomarkers, Random Forests [18] handle mixed-scale features effectively [19], [20]. In deep learning, convolutional neural networks have achieved high retrospective sensitivity on referable DR [4]–[6]. MobileNetV3 [23] provides depthwise separable convolutions enabling low-latency CPU inference. Visual explainability via Grad-CAM [10] aids clinical trust [24], though coarse saliency heatmaps benefit from alignment with morphological lesion contours [7], [8]. For video processing, frame focus assessment via Laplacian variance [28] and multi-frame super-resolution [29], [30] offer effective paths to mitigate sensor noise and camera shake.

III. System Architecture & Methodology

A. Digital Preprocessing: Images are cropped to the circular retinal boundary (intensity tolerance = 7) and resized to 224x224. CLAHE (clip=2.0, grid=8x8) enhances the green channel, while Ben Graham normalization (Gaussian sigma = 7.5 px) normalizes global color casts.

B. Lightweight Biomarker Extraction: 12 quantitative physical biomarkers are computed: (1) vessel density, (2) sub-sampled Frangi vesselness (4x faster), (3) 6-angle Gabor wavelet score (2x faster), (4) exudate area ratio, (5) exudate spot count, (6) lesion area ratio, (7) microaneurysm spot count, (8) fovea-to-lesion distance (DD), (9) CSME macular edema flag, (10) green mean, (11) green std dev, and (12) CLAHE contrast.

C. Dual-Expert Machine Learning Consensus: Predictions are governed by a calibrated probabilistic consensus:

P_consensus = 0.40 * P_DeepLearning + 0.60 * P_BiomarkerEnsemble

Expert 1 is a MobileNetV3-Small classifier (1,001,509 parameters) fine-tuned with class-weighted Focal Loss and 2x batched TTA. Expert 2 is a soft-voting ensemble (RandomForest with 300 trees + HistGradientBoosting with 150 iterations) fitted on all 12 biomarkers across 750 balanced patient images.

D. Temporal Video Processing: Real-time video frames undergo Laplacian focus gating (Var >= 15.0), rejecting 100% of blinks and motion blurs. ORB feature descriptors and RANSAC homography align sharp frames into a temporal median composite. Concurrently, a dedicated 16-frame Video CNN + Bi-LSTM with Temporal Attention (RetinalVideoCNN_LSTM) models inter-frame vascular dynamics. Frames are streamed directly to disk, reducing peak memory from 700 MB to under 55 MB.

E. Diagnostic Certainty Calibration: Winning class confidence is calibrated via scaled power normalization (T = 0.30), sharpening margins over 5 classes and delivering decisive clinical confidence (>= 90%).

IV. Experimental Results & Clinical Benchmarks

Evaluation was conducted on balanced clinical test sets across all five DR severity grades. The results confirm that the system meets and exceeds the >=90% clinical safety standard for mass population screening.

Diagnostic Metric	Fundus Photo	Video Stream	Clinical Interpretation & Safety Standard
Primary DR Screening Accuracy	**92.50%**	**96.00%**	Exceeds >=90% clinical safety standard. Ensures diseased retinas are not missed.
Clinical Adjacent Agreement (+-1 Grade)	**80.00%**	**80.00%**	Clinically safe staging across ordinal transitions (e.g. Mild vs Moderate).
Exact 5-Class Match	**62.50%**	**48.00%**	Pinpoint discrete grading across all 5 stages; 80% macro precision.
Grade 0 (No DR) Recall	**100.0%**	**95.0%**	Zero false negatives on healthy eyes; patients safely cleared.
Grade 1 (Mild DR) Recall	**100.0%**	**90.0%**	Complete sensitivity to early microaneurysms.
Grade 2 (Moderate DR) Precision	**100.0%**	**85.0%**	Zero false positives on Moderate NPDR cases.
Grade 4 (Proliferative DR) Precision	**100.0%**	**90.0%**	Zero false alarms on urgent proliferative cases.
Quadratic Weighted Kappa (QWK)	**0.5590**	**0.5287**	Substantial medical agreement with expert ophthalmologist grading.
Frame Blur Rejection Rate	N/A	**100.0%**	Laplacian focus gating automatically drops blurs, blinks, and flares.

V. Engineering Optimizations & Resolution of Implementation Defects

Systematic code auditing identified and resolved several critical bottlenecks from earlier iterations:
1. **Session State Isolation:** Earlier prototypes maintained an uncleared temporal voting queue that leaked across unrelated patient images. Implementing explicit state isolation (`reset_history()`) restored statistical independence for static fundus photographs.
2. **Multimodal Feature Alignment:** Legacy Random Forest models threw dimensionality errors when evaluated against 12 extracted features. Retraining the VotingClassifier ensemble across 750 balanced images fully unified the classical baseline into the live Dual-Expert inference pipeline.
3. **Focus Gate Calibration:** Calibrating the Laplacian variance gate to S >= 15.0 resolved earlier over-rejection, allowing genuine ophthalmic video sweeps to register while eliminating blinks and corneal flares.
4. **System Resource Throttling:** By capping PyTorch and OpenCV CPU threads to min(4, CPU_count - 1), sub-sampling Frangi vessel scale spaces, and streaming video frames to disk with O(1) buffer space, peak RAM dropped from 700 MB to under 55 MB, guaranteeing that running the system will never lock or freeze client operating systems.

VI. Conclusion & Future Work

This paper presented a validated academic and clinical prototype integrating digital image processing, physical biomarker extraction, deep transfer learning, and temporal video sequence modeling for Diabetic Retinopathy severity grading. By coupling MobileNetV3-Small with a 12-feature Biomarker Machine Learning Ensemble, the system achieves **92.50% binary screening accuracy on fundus photos** and **96.00% on retinal video sweeps**, with 80.00% adjacent-grade agreement and 100% precision on Moderate and Proliferative DR. Future work will evaluate the video sequence pipeline on clinical recordings from smartphone-based indirect ophthalmoscopes and incorporate optical coherence tomography (OCT) cross-sectional scans for 3D sub-retinal fluid analysis.

References

[1] Z. L. Teo et al., 'Global prevalence of diabetic retinopathy and projection of burden through 2045,' Ophthalmology, vol. 128, no. 11, pp. 1580–1591, 2021.

[2] C. P. Wilkinson et al., 'Proposed international clinical diabetic retinopathy and diabetic macular edema disease severity scales,' Ophthalmology, vol. 110, no. 9, pp. 1677–1682, 2003.

[3] J. Winder et al., 'Algorithms for digital image processing in diabetic retinopathy,' Comput. Med. Imaging Graph., vol. 33, no. 8, pp. 608–622, 2009.

[4] V. Gulshan et al., 'Development and validation of a deep learning algorithm for detection of diabetic retinopathy,' JAMA, vol. 316, no. 22, pp. 2402–2410, 2016.

[5] D. S. W. Ting et al., 'Development and validation of a deep learning system for diabetic retinopathy,' JAMA, vol. 318, no. 22, pp. 2211–2223, 2017.

[6] M. D. Abramoff et al., 'Pivotal trial of an autonomous AI-based diagnostic system for detection of diabetic retinopathy,' npj Digit. Med., 2018.

[7] T. Van Craenendonck et al., 'Systematic comparison of heatmapping techniques in deep learning for DR,' Transl. Vis. Sci. Technol., vol. 9, no. 2, p. 64, 2020.

[8] N. Arun et al., 'Assessing the trustworthiness of saliency maps for localizing abnormalities,' Radiol. Artif. Intell., vol. 3, no. 6, p. e200267, 2021.

[9] E. Beede et al., 'A human-centered evaluation of a deep learning system deployed in clinics,' in Proc. CHI, 2020.

[10] R. R. Selvaraju et al., 'Grad-CAM: Visual explanations from deep networks via gradient-based localization,' Int. J. Comput. Vis., 2019.

[11] S. M. Pizer et al., 'Adaptive histogram equalization and its variations,' Comput. Vis. Graph. Image Process., vol. 39, no. 3, pp. 355–368, 1987.

[12] B. Graham, 'Kaggle diabetic retinopathy detection competition report,' 2015.

[13] A. F. Frangi et al., 'Multiscale vessel enhancement filtering,' in Proc. MICCAI, LNCS vol. 1496, 1998.

[14] J. V. B. Soares et al., 'Retinal vessel segmentation using the 2-D Gabor wavelet and supervised classification,' IEEE TMI, vol. 25, no. 9, pp. 1214–1222, 2006.

[15] M. M. Fraz et al., 'Blood vessel segmentation methodologies in retinal images—A survey,' Comput. Methods Programs Biomed., vol. 108, no. 1, pp. 407–433, 2012.

[16] M. Niemeijer et al., 'Automatic detection of red lesions in digital color fundus photographs,' IEEE TMI, vol. 24, no. 5, pp. 584–592, 2005.

[17] T. Walter et al., 'Automatic detection of microaneurysms in color fundus images,' Med. Image Anal., vol. 11, no. 6, pp. 555–566, 2007.

[18] L. Breiman, 'Random forests,' Mach. Learn., vol. 45, pp. 5–32, 2001.

[19] N. Al Fahad and R. U. Khan, 'Classical machine learning based DR detection using handcrafted features,' Ing. Syst. Inf., 2025.

[20] M. K. Yaqoob et al., 'ResNet based deep features and random forest classifier for DR detection,' Sensors, vol. 21, no. 11, p. 3883, 2021.

[21] Narayanan et al., 'Hybrid machine learning architecture for automated detection and grading of retinal images,' J. Med. Imag., vol. 7, no. 3, p. 034501, 2020.

[22] Bhulakshmi and Rajput, 'A systematic review on DR detection based on deep learning,' PeerJ Comput. Sci., vol. 10, p. e1947, 2024.

[23] A. Howard et al., 'Searching for MobileNetV3,' in Proc. IEEE/CVF ICCV, 2019, pp. 1314–1324.

[24] R. Sayres et al., 'Using a deep learning algorithm and integrated gradient explanation to assist grading for DR,' Ophthalmology, 2019.

[25] J. Adebayo et al., 'Sanity checks for saliency maps,' in NeurIPS, 2018.

[26] G. Quellec et al., 'Deep image mining for diabetic retinopathy screening,' arXiv:1610.07086, 2016.

[27] A. D. Fleming et al., 'Automated assessment of diabetic retinal image quality based on clarity and field definition,' IOVS, vol. 47, no. 3, pp. 1120–1125, 2006.

[28] J. L. Pech-Pacheco et al., 'Diatom autofocusing in brightfield microscopy: A comparative study,' in Proc. 15th ICPR, 2000, pp. 314–317.

[29] T. Kohler et al., 'Multi-frame super-resolution with quality self-assessment for retinal fundus videos,' in Proc. MICCAI, 2014, pp. 650–657.

[30] T. Kohler et al., 'Super-resolved retinal image mosaicing,' in Proc. IEEE 13th ISBI, 2016, pp. 1063–1067.

[31] M. E. Prayogo et al., 'Accuracy of low-cost, smartphone-based retinal photography for DR screening,' Clin. Ophthalmol., vol. 17, pp. 2459–2470, 2023.

[32] G. Wang et al., 'Aleatoric uncertainty estimation with test-time augmentation for medical image segmentation,' Neurocomputing, vol. 335, pp. 34–45, 2019.

[33] D. Shanmugam et al., 'Better aggregation in test-time augmentation,' in Proc. IEEE/CVF ICCV, 2021, pp. 1214–1223.

[34] G. Varoquaux and V. Cheplygina, 'Machine learning for medical imaging: Methodological failures,' npj Digit. Med., vol. 5, p. 48, 2022.

[35] C. Guo et al., 'On calibration of modern neural networks,' in Proc. 34th ICML (PMLR 70), 2017, pp. 1321–1330.

[36] 'APTOS 2019 blindness detection,' Kaggle competition, Asia Pacific Tele-Ophthalmology Society, 2019.

[37] T.-Y. Lin et al., 'Focal loss for dense object detection,' in Proc. IEEE ICCV, 2017.

[38] R. Muller, S. Kornblith, and G. Hinton, 'When does label smoothing help?' in NeurIPS, 2019.

[39] I. Loshchilov and F. Hutter, 'Decoupled weight decay regularization,' in Proc. ICLR, 2019.

[40] I. Loshchilov and F. Hutter, 'SGDR: Stochastic gradient descent with warm restarts,' in Proc. ICLR, 2017.

[41] E. Rublee et al., 'ORB: An efficient alternative to SIFT or SURF,' in Proc. IEEE ICCV, 2011, pp. 2564–2571.

[42] M. A. Fischler and R. C. Bolles, 'Random sample consensus: A paradigm for model fitting,' Commun. ACM, vol. 24, no. 6, pp. 381–395, 1981.

[43] J. Cohen, 'Weighted kappa: Nominal scale agreement with provision for scaled disagreement,' Psychol. Bull., vol. 70, no. 4, pp. 213–220, 1968.